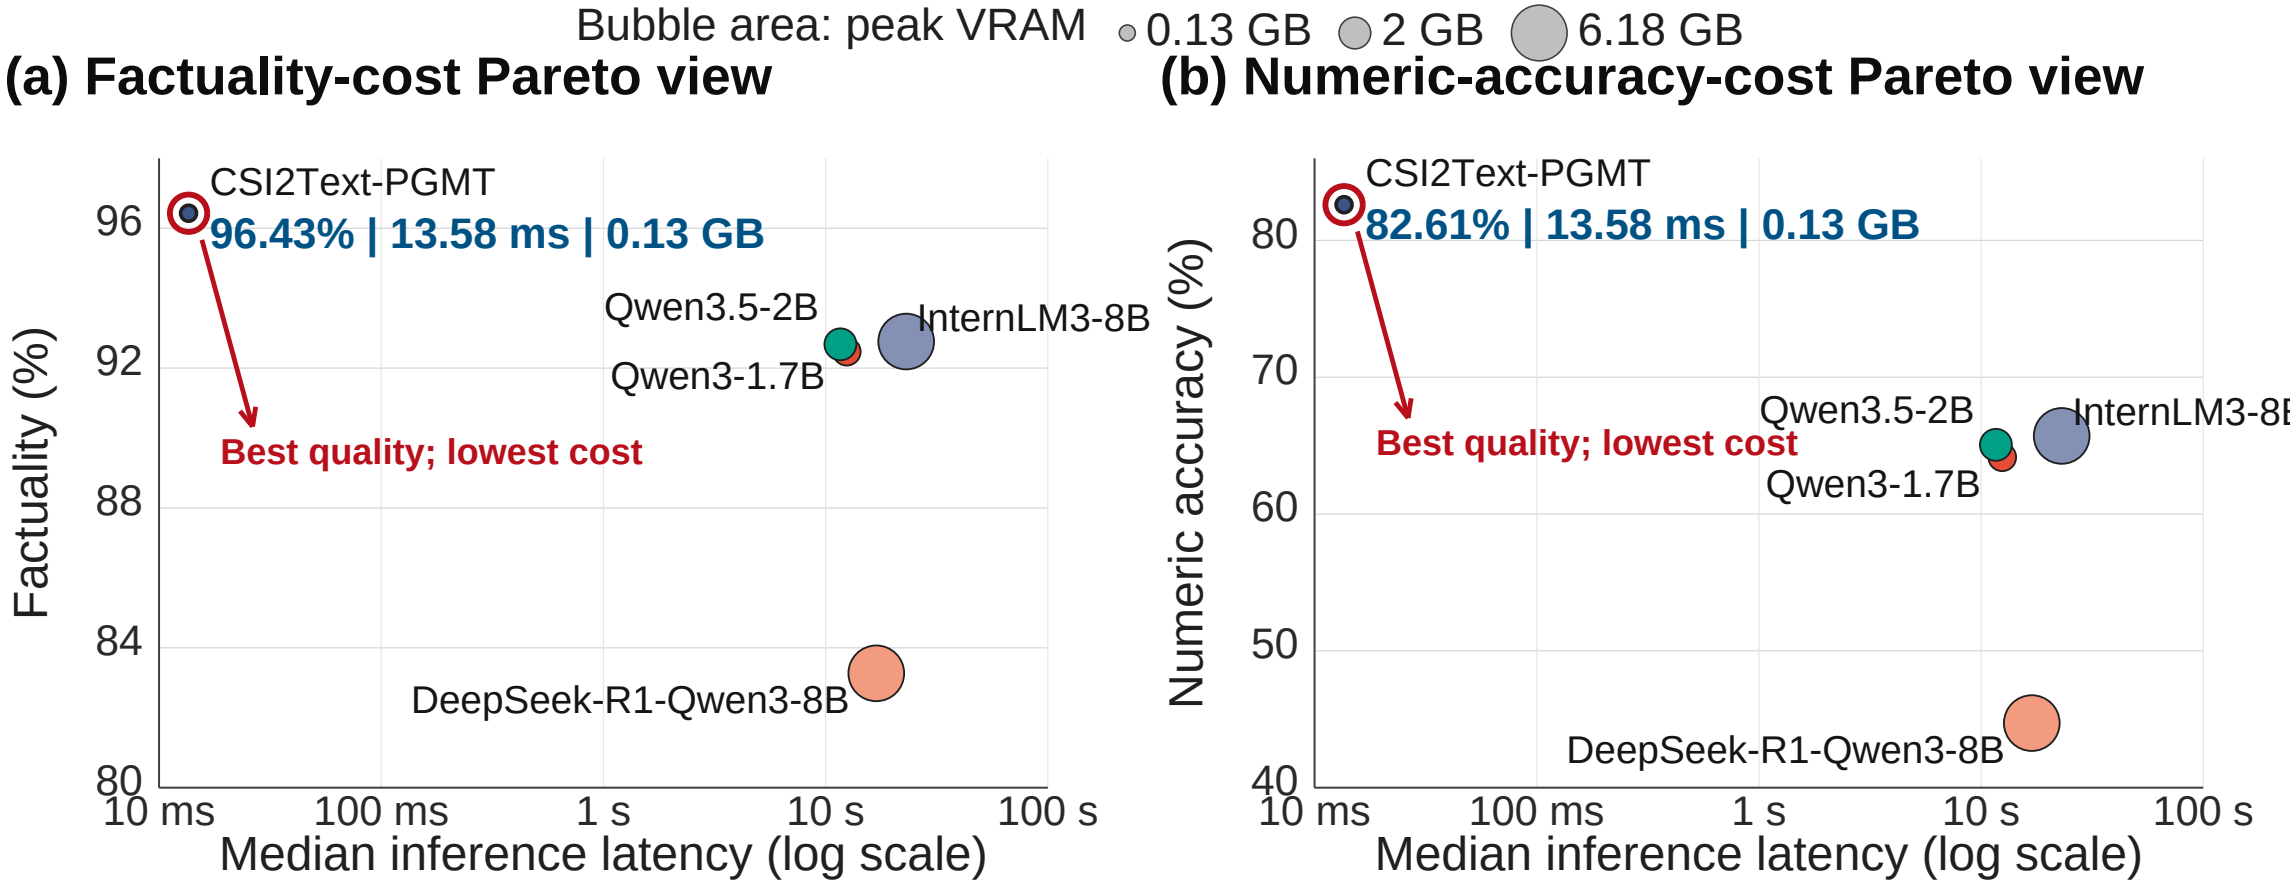
Bubble area: peak VRAM
0.13 GB
2 GB
6.18 GB
(a) Factuality-cost Pareto view
CSI2Text-PGMT
96.43% | 13.58 ms | 0.13 GB
Best quality; lowest cost
Qwen3.5-2B
InternLM3-8B
Qwen3-1.7B
DeepSeek-R1-Qwen3-8B
Factuality (%)
96
92
88
84
80
10 ms
100 ms
1 s
10 s
100 s
Median inference latency (log scale)
(b) Numeric-accuracy-cost Pareto view
CSI2Text-PGMT
82.61% | 13.58 ms | 0.13 GB
Best quality; lowest cost
Qwen3.5-2B
InternLM3-8B
Qwen3-1.7B
DeepSeek-R1-Qwen3-8B
Numeric accuracy (%)
80
70
60
50
40
10 ms
100 ms
1 s
10 s
100 s
Median inference latency (log scale)